Assignment 2 report: The German Traffic Sign Recognition Benchmark — Rev 1

Tianyu Zhang (tz1280@nyu.edu)

Honor Statement: By submitting this work, I certify that, with the exception of LaTeX templates and text whose sources I cite, every keystroke in the answers was typed by me.

Instructor: Prof. Rob Fergus.

1 The environment for this assignment

- OS: Windows 10 x64
- CUDA: 10.0
- Python: 3.7.3
- CUDNN: 7
- pytorch: 1.2.0
- GPU: GeForce GTX 1070 (8+8 GB GPU Memory)
- CPU: i7-8750H

P.S. Estimated time to run one batch on this PC: 5 minutes

2 About the dataset

1. The shape of the sample: [*, 3, 32, 32], which means the sample is RBG with 32 x 32 pixel
2. The training dataset has 35339 samples and the validation dataset has 3870 samples.

3 My framework and the hyper-parameters

I used ensemble models which includes 4 GoogLeNets and a Basic CNN Net. 3 GoogLeNets are trained in the same process (Please see the GoogLeNet-Test-3.ipynb) and the another is trained in another process (Please see the GoogLeNet-Test-2.ipynb). I used the same random seed for the two process of training GoogLeNets. However, the result of the two process are actually different.

How do I choose those GoogLeNets?

In GoogLeNet-Test-3.ipynb, by looking at the log and choose those accuracy are local maximum which means not only that epoch has a big validation accuracy but also the near epochs have relatively larger validation accuracy.

1. GoogLeNet
2. A Basic CNN Net

   I will show the frameworks at the end of this file.

4 How I ensemble these models?

During the process of get the Kaggle files, I got the raw classification scores for each model and take the mean. Then, I use it as the ensemble model result which performs better than any models in the ensemble list.

Please also see the training loss graph and accuracy graph in the end of this doc

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 192, 32, 32]	5,376
BatchNorm2d-2	[-1, 192, 32, 32]	384
ReLU-3	[-1, 192, 32, 32]	0
Conv2d-4	[-1, 64, 32, 32]	12,352
BatchNorm2d-5	[-1, 64, 32, 32]	128
ReLU-6	[-1, 64, 32, 32]	0
Conv2d-7	[-1, 96, 32, 32]	18,528
BatchNorm2d-8	[-1, 96, 32, 32]	192
ReLU-9	[-1, 96, 32, 32]	0
Conv2d-10	[-1, 128, 32, 32]	110,720
BatchNorm2d-11	[-1, 128, 32, 32]	256
ReLU-12	[-1, 128, 32, 32]	0
Conv2d-13	[-1, 16, 32, 32]	3,088
BatchNorm2d-14	[-1, 16, 32, 32]	32
ReLU-15	[-1, 16, 32, 32]	0
Conv2d-16	[-1, 32, 32, 32]	4,640
BatchNorm2d-17	[-1, 32, 32, 32]	64
ReLU-18	[-1, 32, 32, 32]	0
Conv2d-19	[-1, 32, 32, 32]	9,248
BatchNorm2d-20	[-1, 32, 32, 32]	64
ReLU-21	[-1, 32, 32, 32]	0
MaxPool2d-22	[-1, 192, 32, 32]	0
Conv2d-23	[-1, 32, 32, 32]	6,176
BatchNorm2d-24	[-1, 32, 32, 32]	64
ReLU-25	[-1, 32, 32, 32]	0
Inception-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 128, 32, 32]	32,896
BatchNorm2d-28	[-1, 128, 32, 32]	256
ReLU-29	[-1, 128, 32, 32]	0
Conv2d-30	[-1, 128, 32, 32]	32,896
BatchNorm2d-31	[-1, 128, 32, 32]	256
ReLU-32	[-1, 128, 32, 32]	0
Conv2d-33	[-1, 192, 32, 32]	221,376
BatchNorm2d-34	[-1, 192, 32, 32]	384
ReLU-35	[-1, 192, 32, 32]	0
Conv2d-36	[-1, 32, 32, 32]	8,224
BatchNorm2d-37	[-1, 32, 32, 32]	64
ReLU-38	[-1, 32, 32, 32]	0
Conv2d-39	[-1, 96, 32, 32]	27,744
BatchNorm2d-40	[-1, 96, 32, 32]	192
ReLU-41	[-1, 96, 32, 32]	0
Conv2d-42	[-1, 96, 32, 32]	83,040
BatchNorm2d-43	[-1, 96, 32, 32]	192
ReLU-44	[-1, 96, 32, 32]	0
MaxPool2d-45	[-1, 256, 32, 32]	0
Conv2d-46	[-1, 64, 32, 32]	16,448
BatchNorm2d-47	[-1, 64, 32, 32]	128
ReLU-48	[-1, 64, 32, 32]	0
Inception-49	[-1, 480, 32, 32]	0
MaxPool2d-50	[-1, 480, 16, 16]	0
Conv2d-51	[-1, 192, 16, 16]	92,352
BatchNorm2d-52	[-1, 192, 16, 16]	384
ReLU-53	[-1, 192, 16, 16]	0
Conv2d-54	[-1, 96, 16, 16]	46,176
BatchNorm2d-55	[-1, 96, 16, 16]	192
ReLU-56	[-1, 96, 16, 16]	0
Conv2d-57	[-1, 208, 16, 16]	179,920
BatchNorm2d-58	[-1, 208, 16, 16]	416
ReLU-59	[-1, 208, 16, 16]	0
Conv2d-60	[-1, 16, 16, 16]	7,696
BatchNorm2d-61	[-1, 16, 16, 16]	32
ReLU-62	[-1, 16, 16, 16]	0
Conv2d-63	[-1, 48, 16, 16]	6,960
BatchNorm2d-64	[-1, 48, 16, 16]	96
ReLU-65	[-1, 48, 16, 16]	0
Conv2d-66	[-1, 48, 16, 16]	20,784
BatchNorm2d-67	[-1, 48, 16, 16]	96
ReLU-68	[-1, 48, 16, 16]	0
MaxPool2d-69	[-1, 480, 16, 16]	0
Conv2d-70	[-1, 64, 16, 16]	30,784
BatchNorm2d-71	[-1, 64, 16, 16]	128
ReLU-72	[-1, 64, 16, 16]	0
Inception-73	[-1, 512, 16, 16]	0
Conv2d-74	[-1, 160, 16, 16]	82,080
BatchNorm2d-75	[-1, 160, 16, 16]	320
ReLU-76	[-1, 160, 16, 16]	0
Conv2d-77	[-1, 112, 16, 16]	57,456
BatchNorm2d-78	[-1, 112, 16, 16]	224
ReLU-79	[-1, 112, 16, 16]	0
Conv2d-80	[-1, 224, 16, 16]	226,016
BatchNorm2d-81	[-1, 224, 16, 16]	448
ReLU-82	[-1, 224, 16, 16]	0
Conv2d-83	[-1, 24, 16, 16]	12,312
BatchNorm2d-84	[-1, 24, 16, 16]	48
ReLU-85	[-1, 24, 16, 16]	0
Conv2d-86	[-1, 64, 16, 16]	13,888
BatchNorm2d-87	[-1, 64, 16, 16]	128
ReLU-88	[-1, 64, 16, 16]	0
Conv2d-89	[-1, 64, 16, 16]	36,928
BatchNorm2d-90	[-1, 64, 16, 16]	128
ReLU-91	[-1, 64, 16, 16]	0
MaxPool2d-92	[-1, 512, 16, 16]	0
Conv2d-93	[-1, 64, 16, 16]	32,832
BatchNorm2d-94	[-1, 64, 16, 16]	128
ReLU-95	[-1, 64, 16, 16]	0
Inception-96	[-1, 512, 16, 16]	0
Conv2d-97	[-1, 128, 16, 16]	65,664
BatchNorm2d-98	[-1, 128, 16, 16]	256
ReLU-99	[-1, 128, 16, 16]	0
Conv2d-100	[-1, 128, 16, 16]	65,664
BatchNorm2d-101	[-1, 128, 16, 16]	256
ReLU-102	[-1, 128, 16, 16]	0
Conv2d-103	[-1, 256, 16, 16]	295,168
BatchNorm2d-104	[-1, 256, 16, 16]	512
ReLU-105	[-1, 256, 16, 16]	0
Conv2d-106	[-1, 24, 16, 16]	12,312
BatchNorm2d-107	[-1, 24, 16, 16]	48
ReLU-108	[-1, 24, 16, 16]	0
Conv2d-109	[-1, 64, 16, 16]	13,888
BatchNorm2d-110	[-1, 64, 16, 16]	128
ReLU-111	[-1, 64, 16, 16]	0
Conv2d-112	[-1, 64, 16, 16]	36,928
BatchNorm2d-113	[-1, 64, 16, 16]	128
ReLU-114	[-1, 64, 16, 16]	0
MaxPool2d-115	[-1, 512, 16, 16]	0
Conv2d-116	[-1, 64, 16, 16]	32,832
BatchNorm2d-117	[-1, 64, 16, 16]	128
ReLU-118	[-1, 64, 16, 16]	0
Inception-119	[-1, 512, 16, 16]	0
Conv2d-120	[-1, 112, 16, 16]	57,456
BatchNorm2d-121	[-1, 112, 16, 16]	224
ReLU-122	[-1, 112, 16, 16]	0
Conv2d-123	[-1, 144, 16, 16]	73,872
BatchNorm2d-124	[-1, 144, 16, 16]	288
ReLU-125	[-1, 144, 16, 16]	0
Conv2d-126	[-1, 288, 16, 16]	373,536
BatchNorm2d-127	[-1, 288, 16, 16]	576
ReLU-128	[-1, 288, 16, 16]	0
Conv2d-129	[-1, 32, 16, 16]	16,416
BatchNorm2d-130	[-1, 32, 16, 16]	64
ReLU-131	[-1, 32, 16, 16]	0
Conv2d-132	[-1, 64, 16, 16]	18,496
BatchNorm2d-133	[-1, 64, 16, 16]	128
ReLU-134	[-1, 64, 16, 16]	0
Conv2d-135	[-1, 64, 16, 16]	36,928
BatchNorm2d-136	[-1, 64, 16, 16]	128
ReLU-137	[-1, 64, 16, 16]	0
MaxPool2d-138	[-1, 512, 16, 16]	0
Conv2d-139	[-1, 64, 16, 16]	32,832
BatchNorm2d-140	[-1, 64, 16, 16]	128
ReLU-141	[-1, 64, 16, 16]	0
Inception-142	[-1, 528, 16, 16]	0
Conv2d-143	[-1, 256, 16, 16]	135,424
BatchNorm2d-144	[-1, 256, 16, 16]	512
ReLU-145	[-1, 256, 16, 16]	0
Conv2d-146	[-1, 160, 16, 16]	84,640
BatchNorm2d-147	[-1, 160, 16, 16]	320
ReLU-148	[-1, 160, 16, 16]	0
Conv2d-149	[-1, 320, 16, 16]	461,120
BatchNorm2d-150	[-1, 320, 16, 16]	640
ReLU-151	[-1, 320, 16, 16]	0
Conv2d-152	[-1, 32, 16, 16]	16,928
BatchNorm2d-153	[-1, 32, 16, 16]	64
ReLU-154	[-1, 32, 16, 16]	0
Conv2d-155	[-1, 128, 16, 16]	36,992
BatchNorm2d-156	[-1, 128, 16, 16]	256
ReLU-157	[-1, 128, 16, 16]	0
Conv2d-158	[-1, 128, 16, 16]	147,584
BatchNorm2d-159	[-1, 128, 16, 16]	256
ReLU-160	[-1, 128, 16, 16]	0
MaxPool2d-161	[-1, 528, 16, 16]	0
Conv2d-162	[-1, 128, 16, 16]	67,712
BatchNorm2d-163	[-1, 128, 16, 16]	256
ReLU-164	[-1, 128, 16, 16]	0
Inception-165	[-1, 832, 16, 16]	0
MaxPool2d-166	[-1, 832, 8, 8]	0
Conv2d-167	[-1, 256, 8, 8]	213,248
BatchNorm2d-168	[-1, 256, 8, 8]	512
ReLU-169	[-1, 256, 8, 8]	0
Conv2d-170	[-1, 160, 8, 8]	133,280
BatchNorm2d-171	[-1, 160, 8, 8]	320
ReLU-172	[-1, 160, 8, 8]	0
Conv2d-173	[-1, 320, 8, 8]	461,120
BatchNorm2d-174	[-1, 320, 8, 8]	640
ReLU-175	[-1, 320, 8, 8]	0
Conv2d-176	[-1, 32, 8, 8]	26,656
BatchNorm2d-177	[-1, 32, 8, 8]	64
ReLU-178	[-1, 32, 8, 8]	0
Conv2d-179	[-1, 128, 8, 8]	36,992
BatchNorm2d-180	[-1, 128, 8, 8]	256
ReLU-181	[-1, 128, 8, 8]	0
Conv2d-182	[-1, 128, 8, 8]	147,584
BatchNorm2d-183	[-1, 128, 8, 8]	256
ReLU-184	[-1, 128, 8, 8]	0
MaxPool2d-185	[-1, 832, 8, 8]	0
Conv2d-186	[-1, 128, 8, 8]	106,624
BatchNorm2d-187	[-1, 128, 8, 8]	256
ReLU-188	[-1, 128, 8, 8]	0
Inception-189	[-1, 832, 8, 8]	0
Conv2d-190	[-1, 384, 8, 8]	319,872
BatchNorm2d-191	[-1, 384, 8, 8]	768
ReLU-192	[-1, 384, 8, 8]	0
Conv2d-193	[-1, 192, 8, 8]	159,936
BatchNorm2d-194	[-1, 192, 8, 8]	384
ReLU-195	[-1, 192, 8, 8]	0
Conv2d-196	[-1, 384, 8, 8]	663,936
BatchNorm2d-197	[-1, 384, 8, 8]	768
ReLU-198	[-1, 384, 8, 8]	0
Conv2d-199	[-1, 48, 8, 8]	39,984
BatchNorm2d-200	[-1, 48, 8, 8]	96
ReLU-201	[-1, 48, 8, 8]	0
Conv2d-202	[-1, 128, 8, 8]	55,424
BatchNorm2d-203	[-1, 128, 8, 8]	256
ReLU-204	[-1, 128, 8, 8]	0
Conv2d-205	[-1, 128, 8, 8]	147,584
BatchNorm2d-206	[-1, 128, 8, 8]	256
ReLU-207	[-1, 128, 8, 8]	0
MaxPool2d-208	[-1, 832, 8, 8]	0
Conv2d-209	[-1, 128, 8, 8]	106,624
BatchNorm2d-210	[-1, 128, 8, 8]	256
ReLU-211	[-1, 128, 8, 8]	0
Inception-212	[-1, 1024, 8, 8]	0
AvgPool2d-213	[-1, 1024, 1, 1]	0
Dropout2d-214	[-1, 1024, 1, 1]	0
Linear-215	[-1, 43]	44,075

Total params: 6,200,075
Trainable params: 6,200,075
Non-trainable params: 0

Input size (MB): 0.01
Forward/backward pass size (MB): 81.43
Params size (MB): 23.65
Estimated Total Size (MB): 105.09

Figure 1: Model summary for GoogLeNet

```
----------------------------------------------------------------
        Layer (type)               Output Shape         Param #
================================================================
            Conv2d-1          [-1, 100, 30, 30]           2,800
            Conv2d-2          [-1, 150, 12, 12]         240,150
         Dropout2d-3          [-1, 150, 12, 12]               0
            Conv2d-4            [-1, 250, 4, 4]         337,750
         Dropout2d-5            [-1, 250, 4, 4]               0
            Linear-6                  [-1, 200]         200,200
            Linear-7                   [-1, 43]           8,643
================================================================
Total params: 789,543
Trainable params: 789,543
Non-trainable params: 0
----------------------------------------------------------------
Input size (MB): 0.01
Forward/backward pass size (MB): 1.08
Params size (MB): 3.01
Estimated Total Size (MB): 4.10
----------------------------------------------------------------
```

Figure 2: Model summary for A Basic CNN Net

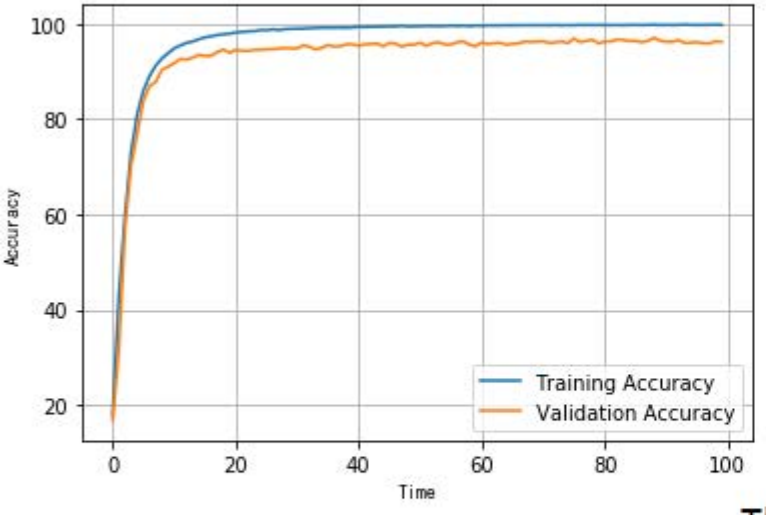

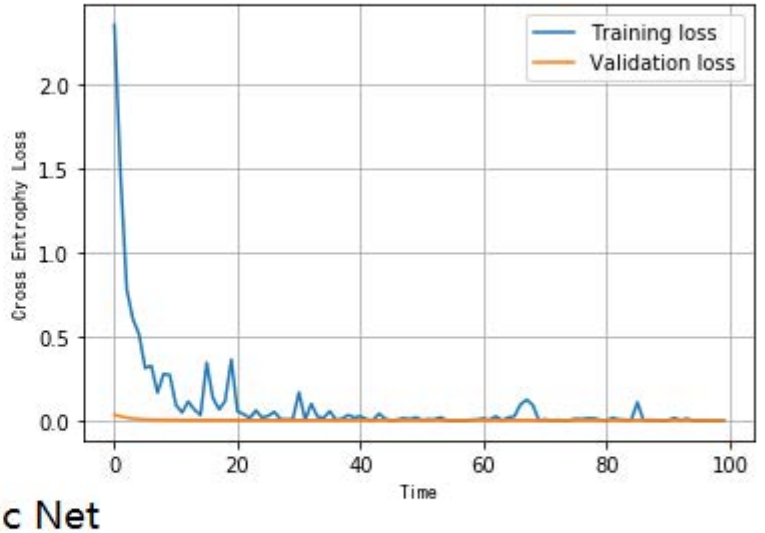

The Basic Net

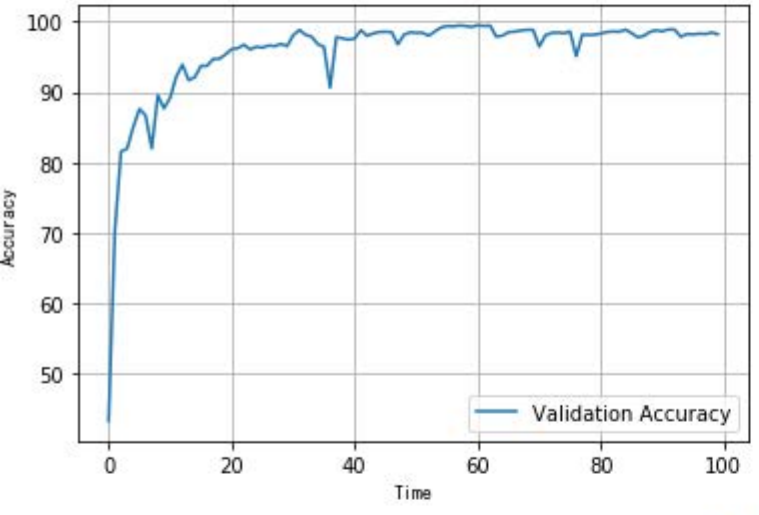

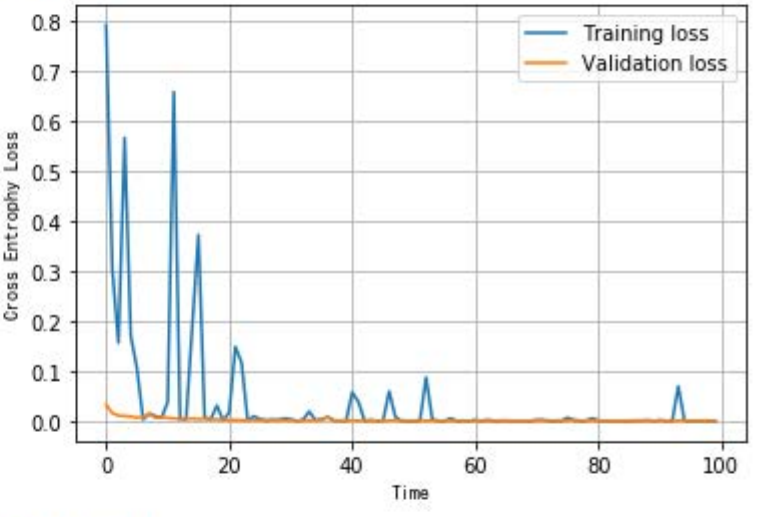

GoogLeNet-test-2

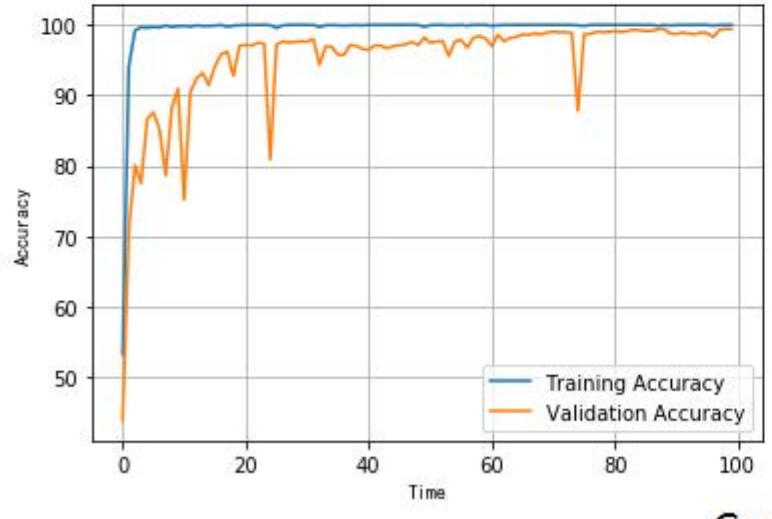

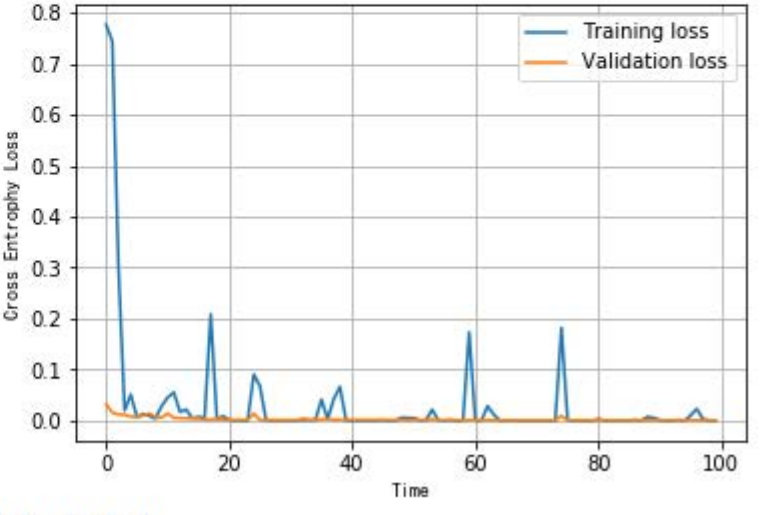

GoogLeNet-test-3